



MLOps Course Final Project - Assessment Criteria

Course Period: January 5, 2026 - March 15, 2026

Final Deadline: March 15, 2026, 24:00 (UTC+1)

Submission Format: GitHub Repository

Group Size: 4 students per group



Project Overview

The final project aims to assess students' ability to **design, implement, and manage an end-to-end MLOps pipeline** using the tools and practices introduced in the course.

Each group will:

- Independently define a **machine learning task**, dataset, and evaluation strategy
 - Implement the project following **MLOps best practices**
 - Track progress and collaboration through **GitHub**
 - Deliver a **final project report** inside the repository
-



Core Technologies Covered

Your project should meaningfully incorporate most of the following:

- **UV** (Python environment & dependency management)
- **Git & GitHub** (collaboration, commits, PRs)
- **Pre-commit & Unit Testing**
- **FastAPI** (model inference service)
- **Docker**

- **MLflow** (experiment tracking & model registry)
 - **Monitoring** (basic metrics, logging, health checks)
 - **CI/CD**
-



Project Timeline & Checkpoints

The project is divided into **4 major checkpoints** in February and March.

✓ Checkpoint 1 — Project Setup & Foundations

Due: February 1, 2026, 24:00

Required Deliverables

- GitHub repository created and shared
- All team members have **write access**
- Python environment managed with **UV**
- Initial project structure (clear, modular, readable)
- Basic data loading and preprocessing logic
- A runnable training script (even a simple baseline model)
- Dependencies correctly tracked (`pyproject.toml` , `uv.lock`)
- Clear README with:
 - Project description
 - Task definition
 - Dataset source
 - Team member roles

Assessed Skills

- Git fundamentals
- Environment reproducibility

- Code organisation
 - Early project planning
-

Checkpoint 2 — Code Quality & Experiment Tracking

 Due: February 15, 2026, 24:00

Required Deliverables

- Pre-commit hooks configured
- Unit tests cover 60%
- Tests runnable locally
- MLflow integrated with log parameters, metrics, model artifacts
- Clear experiment naming and comparison
- Meaningful Git commit history
- README.md updated

Assessed Skills

- Software engineering discipline
 - Testing mindset
 - Reproducible ML experiments
 - Experiment tracking with MLflow
-

Checkpoint 3 — Model Serving & Containerisation

 Due: March 1, 2026, 24:00

Required Deliverables

- FastAPI application for model inference
- Clear API schema (request & response)

- Dockerfile for training and inference
- Application runnable via Docker
- Basic API test
- Inference model loaded from MLflow or saved artifacts
- README.md updated

Assessed Skills

- Model deployment concepts
- API-based ML services
- Containerisation
- Practical usability of the system

Checkpoint 4 — Monitoring, Polish & Final Report

 Due: March 15, 2026, 24:00

Required Deliverables

- Basic monitoring strategy
- Discussion of potential future work
- Final project report (see below)
- Clean, well-documented GitHub repository



Final Project Report in form of README.md

Your report should include:

1. **Project Overview**
2. **Problem Definition & Data**
3. **System Architecture**
4. **MLOps Practices**

5. Monitoring & Reliability

6. Team Collaboration

7. Limitations & Future Work



Only commits made before March 15, 24:00 will be graded.



Build, Deployment & App Demo (Video)

Each group should submit a **demo video (≤ 10 minutes)** showing:

- **Automatic build & deployment triggered by GitHub**
 - CI pipeline, Docker build, app startup
- **How to use the application**
 - FastAPI endpoint(s)
 - Example request and response

Submission:

- Provide a **public or unlisted video link**
 - Include the link in `reports/README.md`
-



Grading Rubric (100 Points)

Category	Points
Project Setup & Git Usage	15
Code Quality & Testing	15
MLflow & Experiment Tracking	15
Model Serving (FastAPI + Docker)	15
Monitoring & Reliability	10
Final Report Quality	15
Bonus	10
Total	100

Teamwork & Individual Contribution

- Git commit history will be reviewed
- Each Pull Request should be reviewed and validated by **at least** one group member.
- All team members should understand **all parts** of the project

|  Unequal contribution may lead to individual score adjustments.

Evaluation Philosophy

This project emphasises:

- **Engineering thinking over model performance**
- **Reproducibility over novelty**
- **Clarity over complexity**

A simple model with excellent MLOps practices can score higher than a complex model with poor engineering.

Good Luck!

If you have questions during the process, **ask us early and often** 😊

yue.li@ext.devinci.fr

linghao.kong@ext.devinci.fr

Good Example

<https://github.com/nielstiben/MLOPS-Project.git>